

Socket programming

Complement for the programming assignment


INFO-0010

Outline

- Socket definition
- Briefing on the Socket API
- A simple example in Java
- Multi-threading and Synchronization
- Example : HTTP protocol
- Debugging tools
- Project overview

What are sockets?

- Interface to network protocol stack
 - Typically to transport layer



What are sockets? (2)


- A socket is an end-point of communication which identifies a local “process” at one end of a communication association
 - A socket is a half association
 - { protocol, local-address, local-port }
- A communication association is identified by two half associations
 - {
 protocol,
 local-address, local-port,
 remote-address, remote-port
}

Communication models

- Datagrams (UDP)
 - Message-oriented
 - Connectionless
 - Unreliable
 - No congestion control
- Connections (TCP)
 - Stream-oriented
 - Requires a connection
 - Reliable (no packet loss, no reordering)
 - Congestion control
- Raw
 - For traffic generation/capture

UDP vs TCP

- Conceptually:



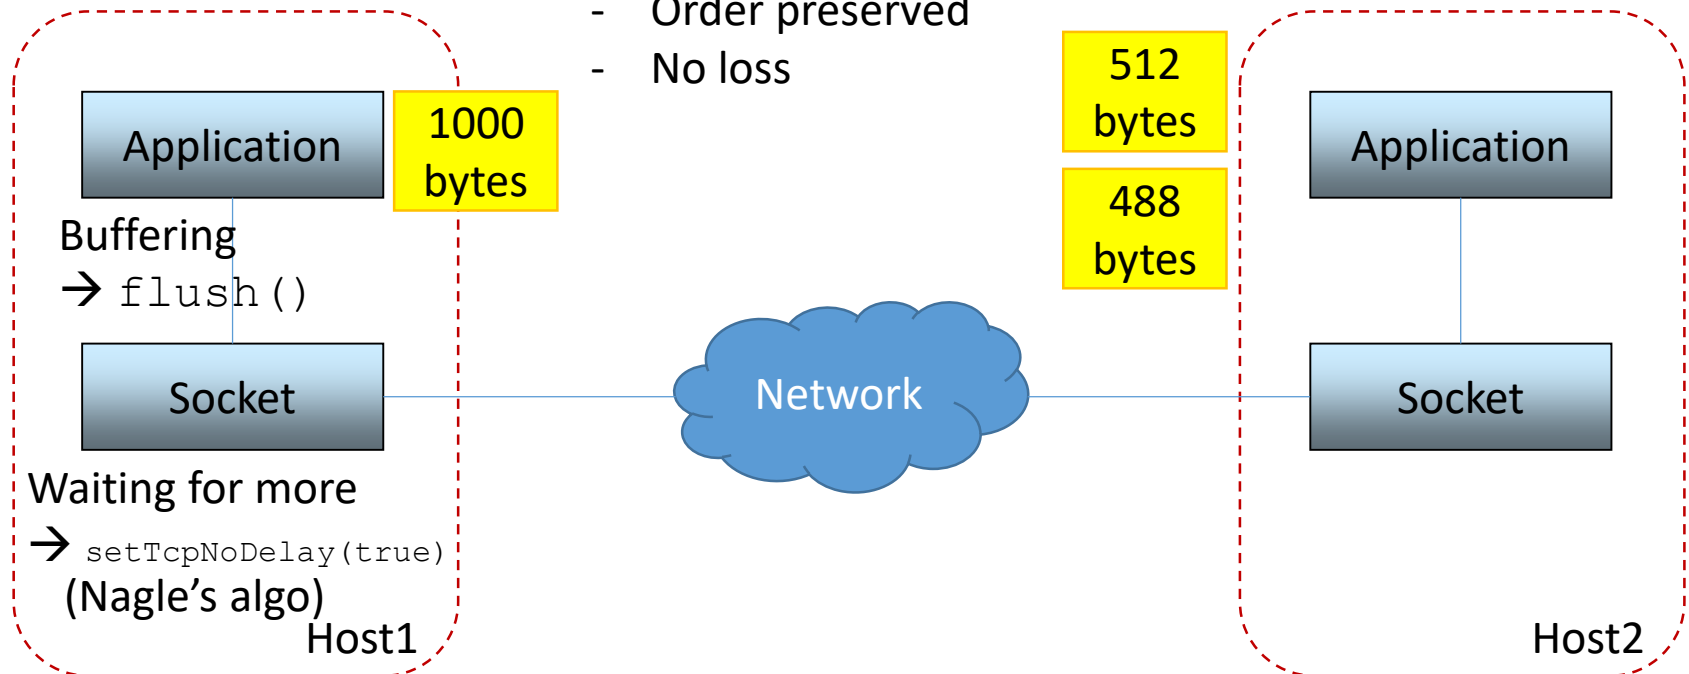
UDP

UDP vs TCP

- Conceptually:

Stream oriented:

- May require multiple reads
- Order preserved
- No loss




TCP

Connections

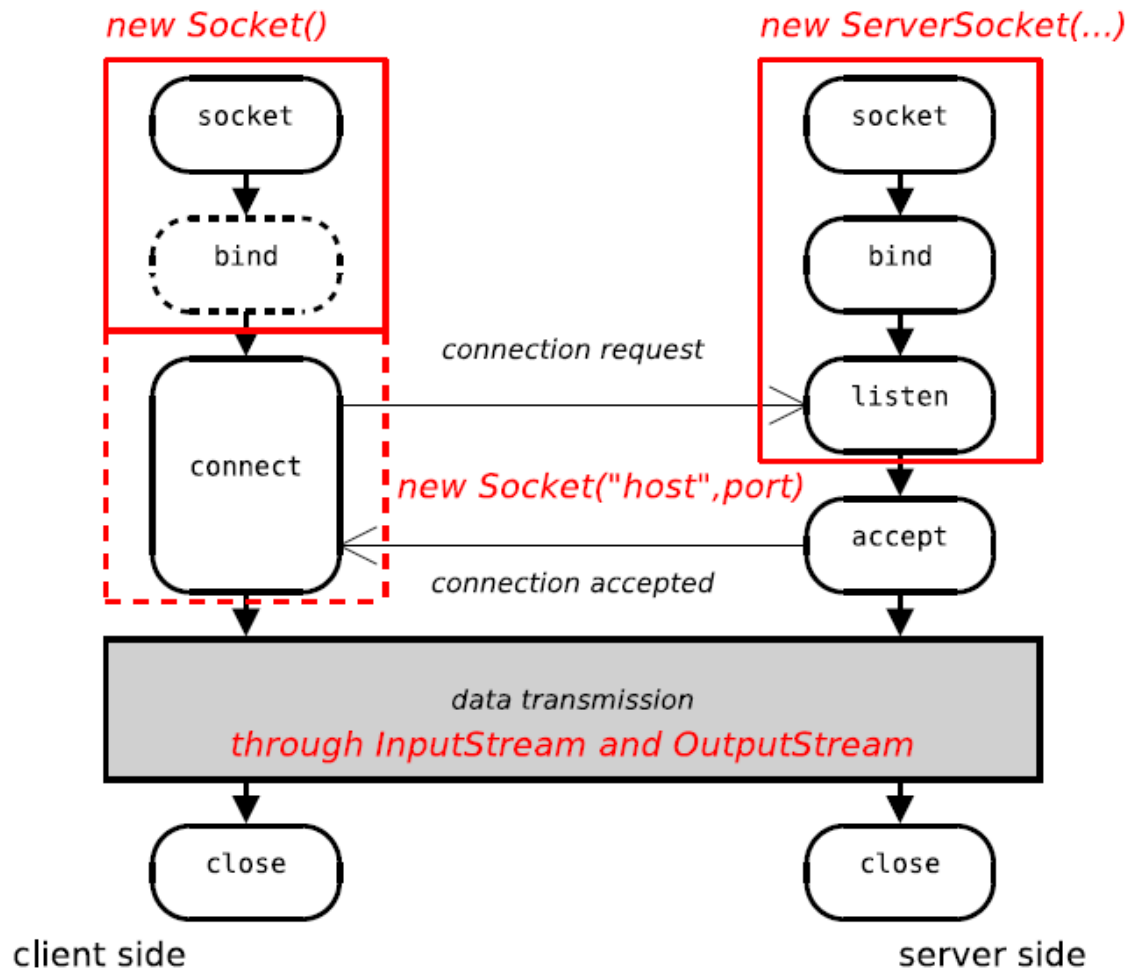
Implemented by TCP

- Reliable stream transfer
 - Chunks read may be different from chunks sent
 - ✓ But streams are identical
- Guarantees delivery and ordering provided connection not broken
- Does congestion control
 - What you have sent to the socket may not have left the box yet!
 - ✓ You can use `out.flush()` to force the writing to the socket
 - ✓ You can use `socket.setTcpNoDelay(true)` to disable Nagle's algorithm
- Programmer must check how many bytes should be read
 - Convention or application protocol header
 - InputStream can provide `readFully()`

Sockets' life cycle (syscalls)




Socket's life cycle (Java)



Simple example : client/server chat

- `telnet` towards a well-known server, from each client
- The server duplicates each incoming message to all the clients.



The Client Side – a bot

```
import java.io.*;
import java.net.*;
```

```
class Bot {
```

```
    public static void main (String argv [ ] ) throws Exception {
```

```
        Socket s = new Socket ( "localhost", 8086 );
```

```
        OutputStream out = s.getOutputStream ( );
```

```
        InputStream in = s.getInputStream ( );
```

```
        byte msg[ ] = new byte [64];
```

```
        out.write ( "ANnA joined the channel" .getBytes ( ) );
```

```
        while ( true ) {
```

```
            if (in.read(msg) <= 0) break;
```

```
            if (new String(msg).startsWith("ANnA") )
```

```
                out.write( "ANnA feels fine , thanks.\n" . getBytes ( ) );
```

```
        }
```

```
        s.close ( );
```

```
    }
```

```
}
```

main() should always catch exceptions.
Would only make the code harder to read in this example.

Use explicit variable names.

We mainly focus on the class names in this example.

The Server Side – Incoming Connection

```
class Server {  
    public static void main ( String argv [ ] ) throws Exception {  
        ServerSocket ss = new ServerSocket (8086) ;  
        while ( true ) {  
            Socket ts = ss.accept () ;  
            OutputStream out = ts.getOutputStream() ;  
            InputStream in = ts.getInputStream() ;  
            out.write("Hello, this is the echo server". getBytes()) ;  
            byte msg [ ] = new byte [64] ;  
  
            while ( true ) {  
                int len= in.read(msg);    // get bytes (max 64)  
                if (len <=0) break ;        // connection closed by peer ?  
                out.write(msg,0,len);      // send them away .  
                out.flush();               // don't wait for more .  
            }  
            ts .close ();  
        }  
    }  
}
```

What if multiple clients connect simultaneously?

The Server Side – multithreading

```
class Server {  
    public static void main ( String argv [ ] ) throws Exception {  
        ServerSocket ss = new ServerSocket (8086) ;  
        while ( true ) {  
            Socket ts = ss.accept () ;  
            Worker w = new Worker(ts);  
            w.start();           //Worker extends Thread  
        }  
    }  
}
```

- We spawn a *thread* every time a connection arrives
- That fresh-new thread will deal with the new client.
- And the main thread can return to the welcoming of incoming clients.

The Server Side – defining a thread

```
class Worker extends Thread {  
    Socket s ;  
    Worker ( Socket _s ) { s = _s ; }  
    @Override  
    public void run ( ) {  
        try {  
            OutputStream out = s.getOutputStream() ;  
            InputStream in = s.getInputStream() ;  
            out.write("Hello, this is the echo server". getBytes() ) ;  
            byte msg [ ] = new byte [64] ;  
  
            while ( true ) {  
                int len= in.read(msg);           // get bytes (max 64)  
                if (len <=0) break ;              // connection closed by peer ?  
                out.write(msg,0,len);             // send them away .  
                out.flush();                      // don't wait for more .  
            }  
            s .close ();                          //acknowledge end of connection  
        } catch ( Exception any ) {  
            System.err.println("worker died " + any) ;  
        }  
    }  
}
```

Extending the "Thread" class

- One way to create thread in Java
 1. Define class that '*extends*' class called **java.lang.Thread**
 2. Create object of that class
- Must override the **public void run ();** method and state so by adding **@Override** in the method signature to prevent a warning

Implementing the “Runnable” interface

- Another way to create thread in Java
 1. Define class that *implements* interface called **java.lang.Runnable**
 2. Create object of class Thread, passing object of above class to Thread constructor
- N.B. an *interface* is like a class with method signatures but no method implementations
 - “implementing” interface is like extending class
 - But, in addition, must provide implementations of all method signatures in specified interface
 - In Runnable there’s only one signature
 - **public void run ();**

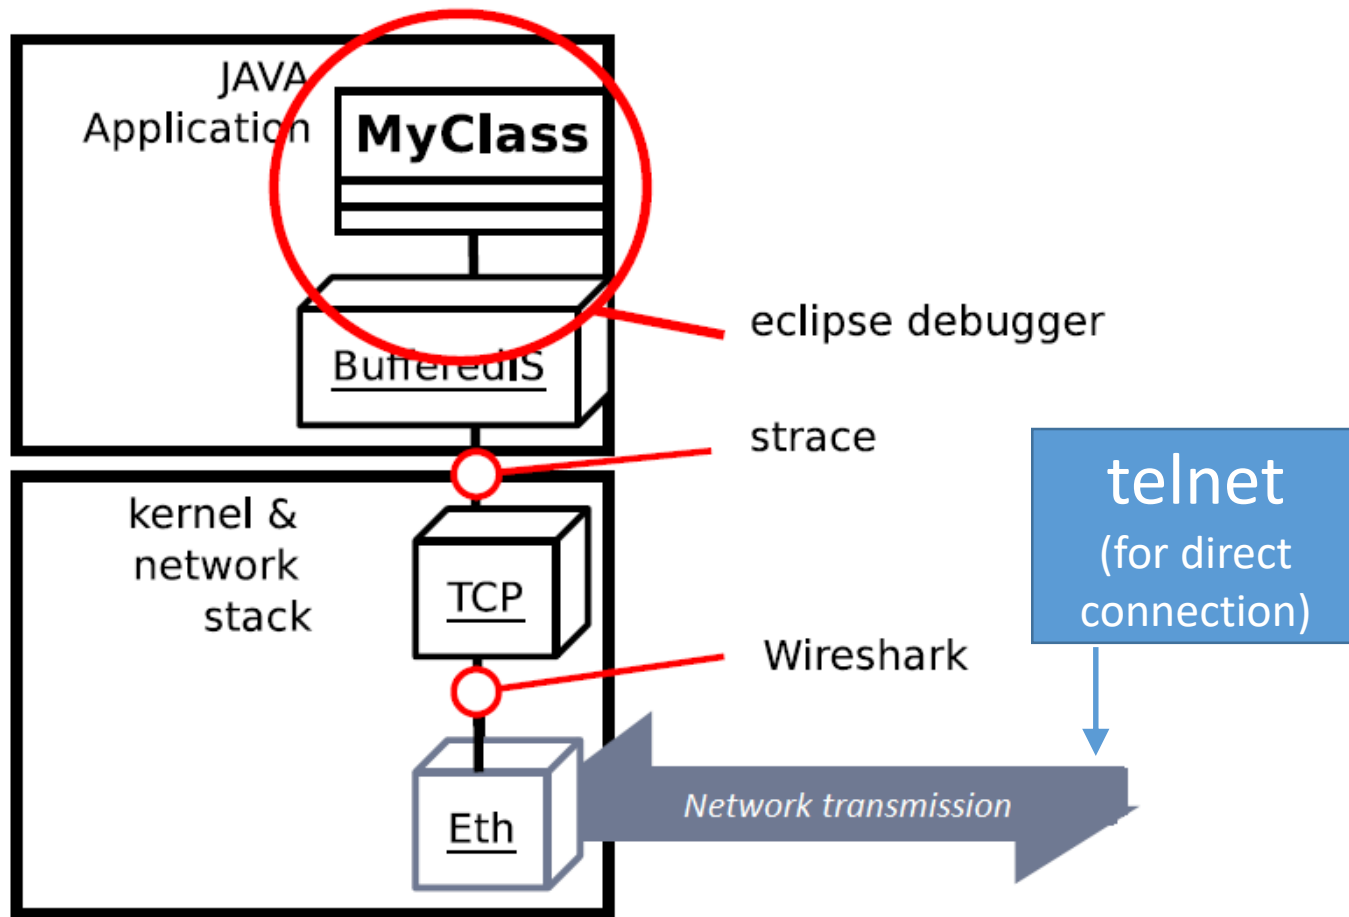
Shared objects

- What if some objects need to be manipulated by different threads?
- For instance, we could keep a list (e.g. `ArrayList`) of all the `OutputStream` and turn the "echo" server into a real chat server.
- Multiple threads will use elements of the list simultaneously
 - Execution is **concurrent** and **non-atomic**
 - Consistency is thus not ensured
- Solution: **only one thread at a time** can use the elements in the list
 - Deem the sending phase a **critical section**
 - Implement **mutual exclusion** over critical section – i.e. prevent multiple threads from entering at once

The Server Side – shared objects

```
// all is just an ArrayList where each OutputStream is .add() ed
// out is the OutputStream corresponding to the Socket from which we are receiving
void send ( byte msg [ ] , int len ) throws Exception {
    synchronized ( all ) {
        for ( Enumeration e = all.elements ( ) ; e.hasMoreElements() ; ) {
            OutputStream o = (OutputStream ) e.nextElement() ;
            if ( o != out ) {
                o.write(msg,0, len ) ;    // send them away .
                o.flush();                // don't wait for more .
            }
        }
    }
}
```

A few debugging tools



Each bug has its proper catcher. So, use `catch` and `printStackTrace(...)` wisely!

strace -enetwork -f java Server

Pid

4199



```
[pid 4199] setsockopt(5, SOL_SOCKET, SO_REUSEADDR, [1], 4) = 0
[pid 4199] bind(5, {sa_family=AF_INET6, sin6_port=htons(8086), inet_pton(AF_INET6, ":::", &sin6_addr), sin6_flowinfo=0, sin6_scope_id=0}, 28) = 0
[pid 4199] listen(5, 50) = 0
[pid 4199] accept(5, {sa_family=AF_INET6, sin6_port=htons(3764), inet_pton(AF_INET6, "::ffff:127.0.0.1", &sin6_addr), sin6_flowinfo=0, sin6_scope_id=0}, [28]) = 6
Process 4211 attached
[pid 4211] sendto(-1251980320, umovestr: Input/output error 0xc, 3086176244, 0, ptrace: umoven: Input/output error {...}, 3042985144) = 0
[pid 4199] accept(5, <unfinished ...>)
[pid 4211] send(6, "Hello, this is the echo server", 30, 0) = 30
[pid 4211] recv(6, <unfinished ...>)
[pid 4199] <... accept resumed> {sa_family=AF_INET6, sin6_port=htons(3765), inet_pton(AF_INET6, "::ffff:127.0.0.1", &sin6_addr), sin6_flowinfo=0, sin6_scope_id=0}, [28]) = 7
Process 4215 attached
[pid 4215] sendto(-1252312096, umovestr: Input/output error 0xc, 3086176244, 0, ptrace: umoven: Input/output error {...}, 3042853368) = 0
[pid 4199] accept(5, <unfinished ...>)
[pid 4215] send(7, "Hello, this is the echo server", 30, 0) = 30
[pid 4215] recv(7, [
```

System calls

strace -enetwork -f java Server

Pid 4199

Pid 4211

Socket N° → 5 = new ServerSocket(8086)

```
[pid 4199] setsockopt(5, SOL_SOCKET, SO_REUSEADDR, [1], 4) = 0
[pid 4199] bind(5, {sa_family=AF_INET6, sin6_port=htons(8086), inet_pton(AF_INET6, ":::", &sin6_addr), sin6_flowinfo=0, sin6_scope_id=0}, 28) = 0
[pid 4199] listen(5, 50) = 0
[pid 4199] accept(5, {sa_family=AF_INET6, sin6_port=htons(3764), inet_pton(AF_INET6, "::ffff:127.0.0.1", &sin6_addr), sin6_flowinfo=0, sin6_scope_id=0}, [28]) = 6
Process 4211 attached
[pid 4211] sendto(-1251980320, umovestr: Input/output error 0xc, 3086176244, 0, ptrace: umoven: Input/output error {...}, 3042385144) = 0
[pid 4199] accept(5, <unfinished ...>)
[pid 4211] send(6, "Hello, this is the echo server", 30, 0) = 30
[pid 4211] recv(6, <unfinished ...>)
[pid 4199] <... accept resumed> {sa_family=AF_INET6, sin6_port=htons(3765), inet_pton(AF_INET6, "::ffff:127.0.0.1", &sin6_addr), sin6_flowinfo=0, sin6_scope_id=0}, [28]) = 7
Process 4215 attached
[pid 4215] sendto(-1252312096, umovestr: Input/output error 0xc, 3086176244, 0, ptrace: umoven: Input/output error {...}, 3042653368) = 0
[pid 4199] accept(5, <unfinished ...>)
[pid 4215] send(7, "Hello, this is the echo server", 30, 0) = 30
[pid 4215] recv(7, [
```

bytes to send

bytes handled by TCP

Some command lines

- (examples are better commented on the web).
- `javac t1.java` to compile
- `java Server` to launch
- `telnet localhost 8086` to test
- `strace -e trace=network -f java Server` to track system calls issued by your program
- `netstat -tlp` to list server sockets
- `netstat -tcp` to list running connections.
- `curl -I http://www.perdu.com` shows the headers for a HEAD method
- `curl http://www.perdu.com` dumps URL content on standard output.